

Name Entity Recognition (NER)

SAMPLE HUGGINGFACE MODELS
WITH MULTINERD-DATA IN
ENGLISH.

Ousted **WeWork** founder **Adam Neumann** lists his **Manhattan** penthouse for **\$37.5 million**

[organization] [person] [location] [monetary value]



Image source: <https://monkeylearn.com/blog/named-entity-recognition/>

Name Entity Recognition (NER)

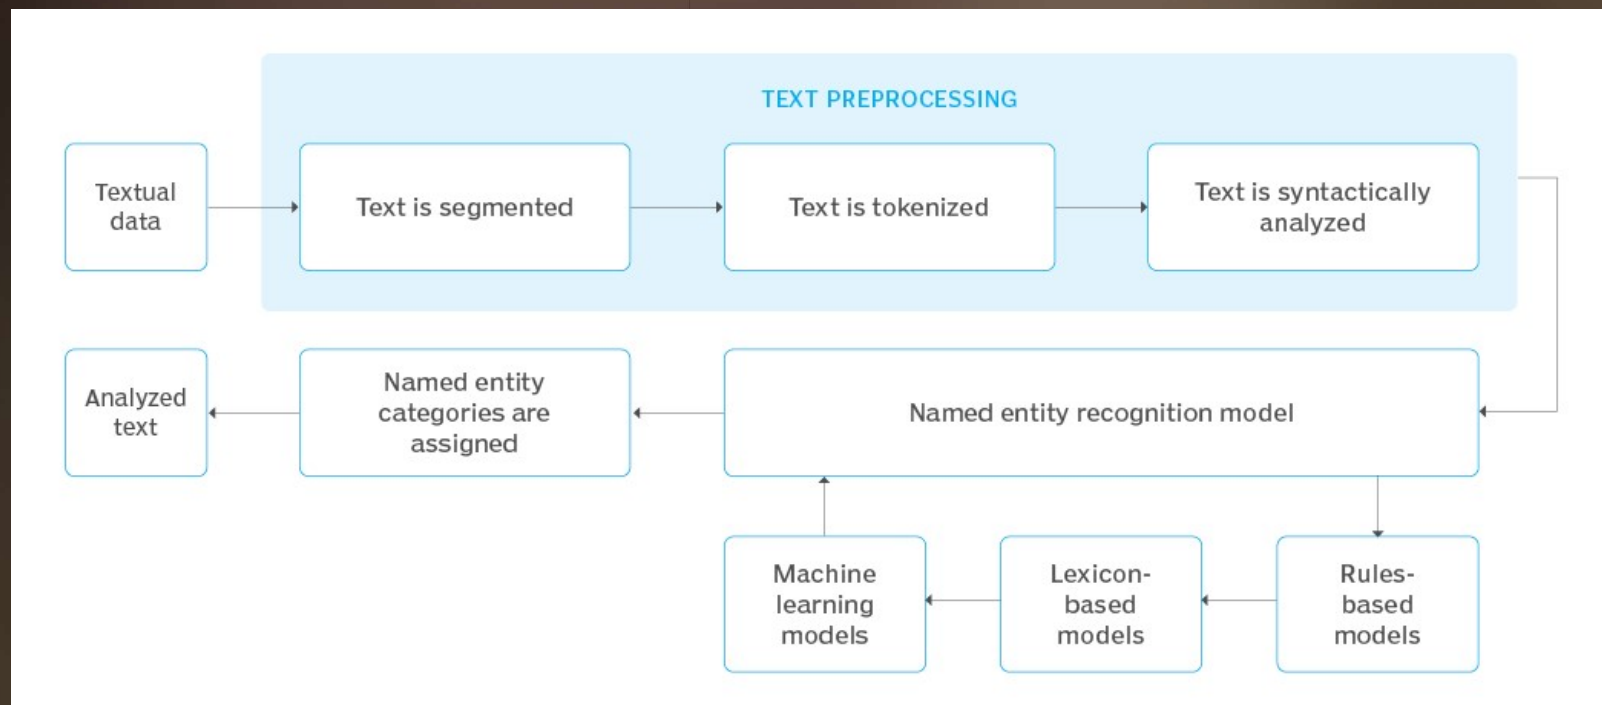
Named entity recognition (NER) – also called entity identification or entity extraction – is a natural language processing (NLP) technique that automatically identifies named entities in a text and classifies them into predefined categories. Entities can be names of people, organizations, locations, times, quantities, monetary values, percentages, and more.

Ousted **WeWork** founder **Adam Neumann** lists his **Manhattan** penthouse for **\$37.5 million**

[organization] [person] [location] [monetary value]

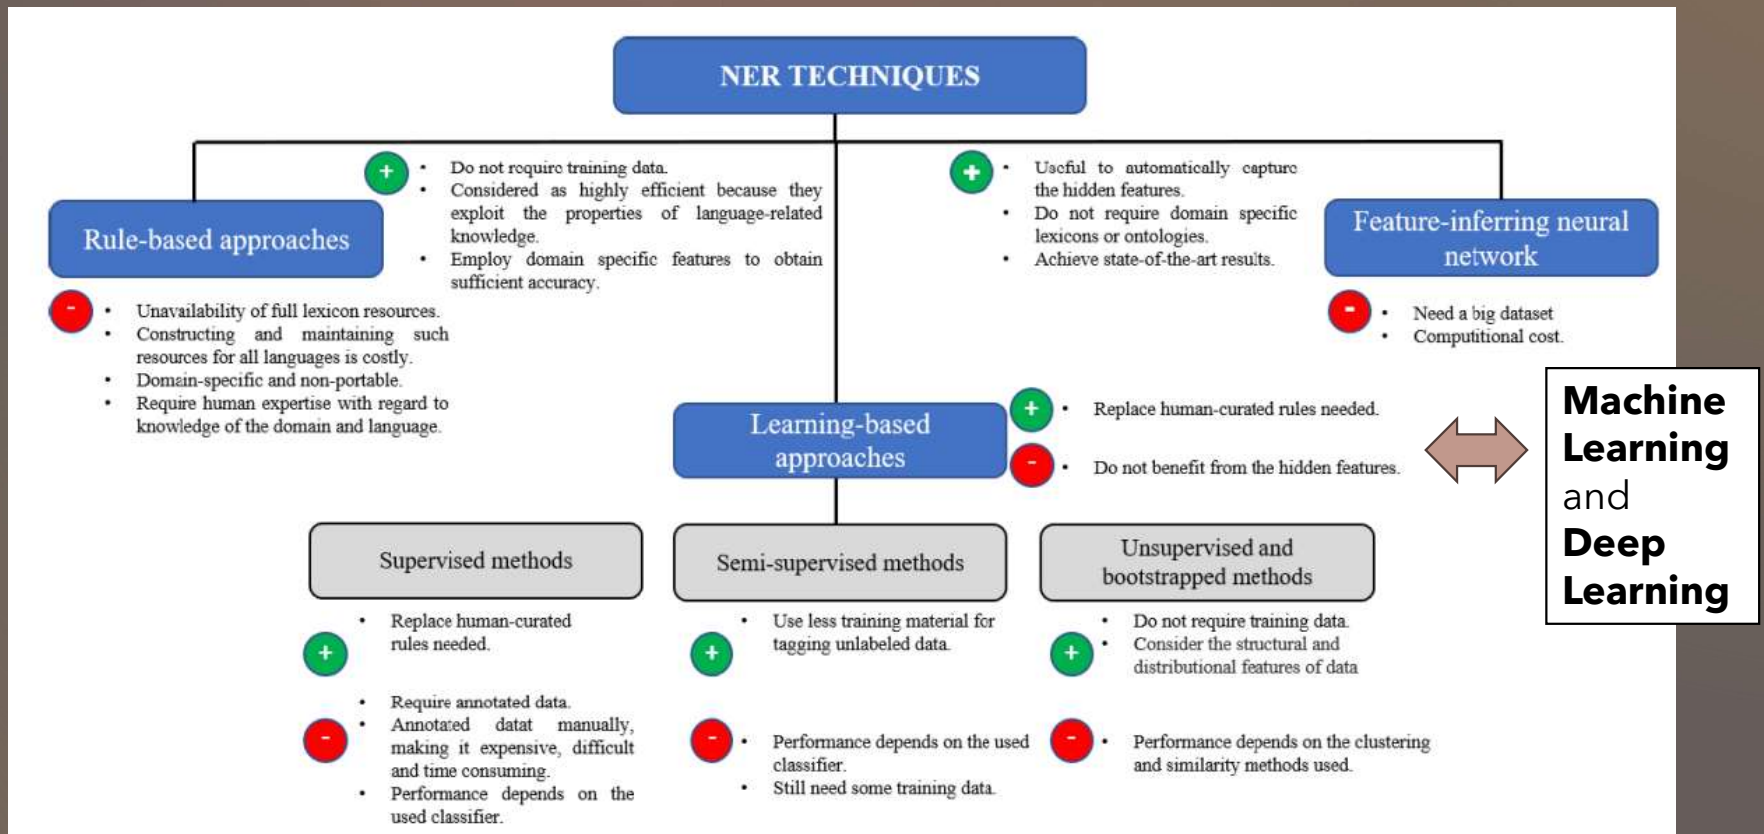
Name Entity Recognition (NER)

-- Process --



Source: <https://www.techtarget.com/whatis/definition/named-entity-recognition-NER>


Current Methods



Source: <https://ieeexplore.ieee.org/ielx7/6287639/8948470/08999622.pdf>

Current Methods

-- Deep Learning --



Source: <https://medium.com/intro-to-artificial-intelligence/entity-extraction-using-deep-learning-8014acac6bb8>

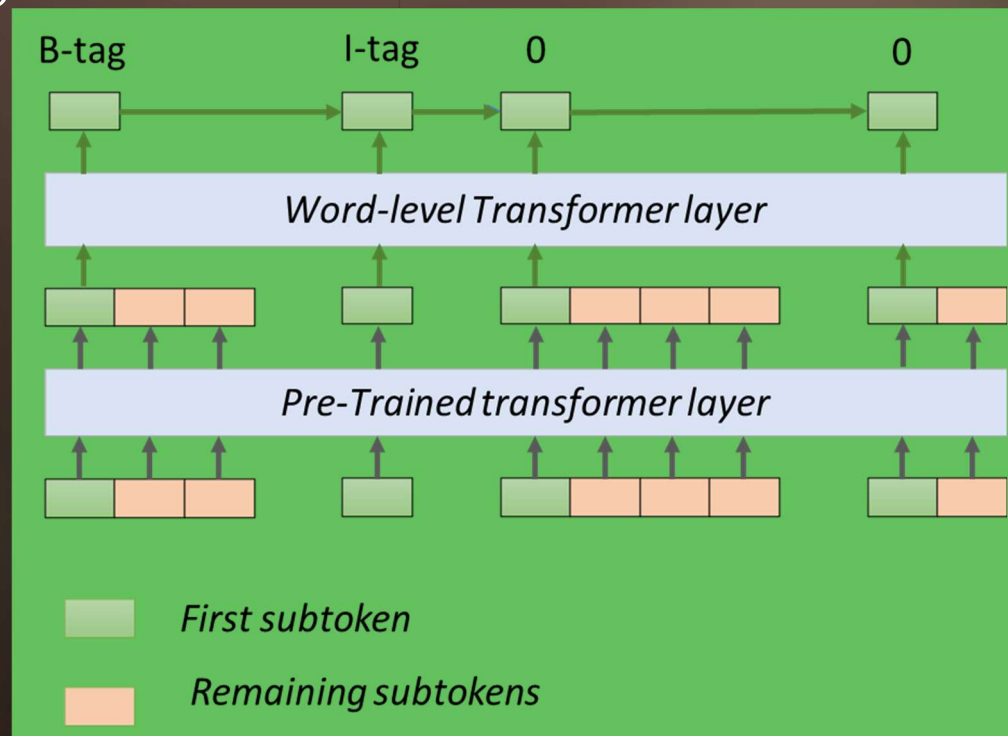
Name Entity Recognition (NER)

-- Deep Learning --

- There are many pre-trained models for this task. For example, you can see the previous paper here (<https://aclanthology.org/2020.coling-main.334.pdf>, <https://pahulpreet86.github.io/name-entity-recognition-pre-trained-models-review/> and <https://pahulpreet86.github.io/name-entity-recognition-pre-trained-models-review/>).
- In this work, I used the Huggingface pre-trained model, namely Distillbert. This is because it is a small model with six deep learning layers and is suitable to run on my machine. You can try with other models such as Roberta and Bert-base. You can find it with the NER task [here](#).

Name Entity Recognition (NER)

-- HuggingFace--



Name Entity Recognition (NER)

-- Distillbert Pre-trained Model --

```
=====
Layer (type:depth-idx)
=====
DistilBertForTokenClassification
├─DistilBertModel: 1-1
│   └─Embeddings: 2-1
│       ├──Embedding: 3-1
│       ├──Embedding: 3-2
│       ├──LayerNorm: 3-3
│       └─Dropout: 3-4
│   └─Transformer: 2-2
│       └─ModuleList: 3-5
├─Dropout: 1-2
└─Linear: 1-3
=====
```


This Repository works

- Scenario A: Select a model from huggingface and run the MultiNERD data for English. Fine-tune the parameters with suitable metrics to get the best results.
- Scenario B: Similar to scenario A. However, the entity types belong to one of the following five entity types: PERSON(PER), ORGANIZATION(ORG), LOCATION(LOC), DISEASES(DIS), ANIMAL(ANIM).

This Repository works

- You can find out the results of SystemA at [SystemADistilbert.ipynb](#)
- Testing the SystemA with an example sentence ([SystemADistilbert_Testing.ipynb](#)).
- You can find out the results of SystemB at [SystemBDistilbert.ipynb](#)
- Testing the SystemB with an example sentence ([SystemBDistilbert_Testing.ipynb](#)).

This Repository works

-- Pre-processing data --

- I only loaded the data and followed the instructions from Huggingface.
- I selected small samples for all training, validation and test sets in two systems because of the limitation of my Machine. You can see it in the codes.
- For SystemB, I only kept the entity types which belong to one of the following five entity types:
**PERSON(PER), ORGANIZATION(ORG),
LOCATION(LOC), DISEASES(DIS), ANIMAL(ANIM).**

A

```
['B-ANIM',  
'B-BIO',  
'B-CEL',  
'B-DIS',  
'B-EVE',  
'B-FOOD',  
'B-INST',  
'B-LOC',  
'B-MEDIA',  
'B-MYTH',  
'B-ORG',  
'B-PER',  
'B-PLANT',  
'B-TIME',  
'B-VEHI',  
'I-ANIM',  
'I-BIO',  
'I-CEL',  
'I-DIS',  
'I-EVE',  
'I-FOOD',  
'I-INST',  
'I-LOC',  
'I-MEDIA',  
'I-MYTH',  
'I-ORG',  
'I-PER',  
'I-PLANT',  
'I-TIME',  
'I-VEHI',  
'O']
```

B

```
['B-ANIM',  
'B-DIS',  
'B-LOC',  
'B-ORG',  
'B-PER',  
'I-ANIM',  
'I-DIS',  
'I-LOC',  
'I-ORG',  
'I-PER',  
'O']
```

This Repository works

-- Pre-processing data --

- For two scenarios, I used the same configuration to compare the outcomes of the two systems.
 - Metrics: F1 score
 - Learning rate: $5e-5 = 0.00005$
 - Epochs: 10

```
training_args = TrainingArguments(  
    output_dir="./results",  
    evaluation_strategy="steps",  
    eval_steps=200,  
    save_steps=200,  
    num_train_epochs=10,  
    per_device_train_batch_size=32,  
    per_device_eval_batch_size=32,  
    logging_steps=100,  
    learning_rate=5e-5,  
    load_best_model_at_end=True,  
    metric_for_best_model="f1",  
)
```

This Repository works

-- Results -- SystemA

Training set

	precision	recall	f1-score	support
ANIM	0.99	0.99	0.99	359
BIO	0.67	0.40	0.50	5
CEL	1.00	1.00	1.00	45
DIS	1.00	1.00	1.00	228
EVE	1.00	1.00	1.00	61
FOOD	1.00	1.00	1.00	217
INST	1.00	0.91	0.95	11
LOC	1.00	1.00	1.00	1442
MEDIA	0.99	1.00	1.00	151
MYTH	1.00	1.00	1.00	13
ORG	1.00	1.00	1.00	654
PER	1.00	1.00	1.00	1515
PLANT	0.99	1.00	0.99	217
TIME	1.00	0.99	0.99	70
VEHI	1.00	1.00	1.00	14
micro avg	1.00	1.00	1.00	5002
macro avg	0.98	0.95	0.96	5002
weighted avg	1.00	1.00	1.00	5002

Validation set

	precision	recall	f1-score	support
ANIM	0.48	0.46	0.47	151
BIO	0.00	0.00	0.00	4
CEL	0.65	0.58	0.61	19
DIS	0.56	0.58	0.57	250
EVE	0.84	0.94	0.88	49
FOOD	0.42	0.38	0.40	329
INST	0.40	0.40	0.40	5
LOC	0.97	0.97	0.97	1195
MEDIA	0.90	0.92	0.91	142
MYTH	1.00	0.33	0.50	9
ORG	0.89	0.92	0.90	404
PER	0.97	0.98	0.98	1122
PLANT	0.40	0.39	0.39	169
TIME	0.56	0.33	0.42	42
VEHI	0.61	0.69	0.65	16
micro avg	0.84	0.83	0.84	3906
macro avg	0.64	0.59	0.60	3906
weighted avg	0.83	0.83	0.83	3906

Test set

	precision	recall	f1-score	support
ANIM	0.51	0.58	0.54	246
BIO	0.00	0.00	0.00	1
CEL	0.00	0.00	0.00	3
DIS	0.63	0.58	0.60	126
EVE	0.83	0.86	0.84	50
FOOD	0.38	0.32	0.34	85
INST	0.00	0.00	0.00	3
LOC	0.97	0.97	0.97	1748
MEDIA	0.86	0.97	0.91	65
MYTH	1.00	0.78	0.88	9
ORG	0.87	0.93	0.90	491
PER	0.97	0.99	0.98	841
PLANT	0.47	0.54	0.50	150
TIME	0.83	0.53	0.65	45
VEHI	0.25	0.60	0.35	5
micro avg	0.87	0.89	0.88	3868
macro avg	0.57	0.58	0.56	3868
weighted avg	0.88	0.89	0.88	3868

This Repository works

-- Results -- SystemB

Training set

	precision	recall	f1-score	support
ANIM	0.97	0.95	0.96	359
DIS	0.92	0.88	0.90	228
LOC	1.00	1.00	1.00	1442
ORG	1.00	1.00	1.00	654
PER	1.00	1.00	1.00	1515
micro avg	0.99	0.99	0.99	4198
macro avg	0.98	0.97	0.97	4198
weighted avg	0.99	0.99	0.99	4198

Validation set

	precision	recall	f1-score	support
ANIM	0.53	0.38	0.44	151
DIS	0.56	0.43	0.49	250
LOC	0.96	0.98	0.97	1195
ORG	0.91	0.88	0.90	404
PER	0.98	0.98	0.98	1122
micro avg	0.92	0.89	0.91	3122
macro avg	0.79	0.73	0.75	3122
weighted avg	0.91	0.89	0.90	3122

Test set

	precision	recall	f1-score	support
ANIM	0.49	0.50	0.49	246
DIS	0.70	0.46	0.56	126
LOC	0.96	0.97	0.96	1748
ORG	0.87	0.89	0.88	491
PER	0.96	0.98	0.97	841
micro avg	0.91	0.91	0.91	3452
macro avg	0.80	0.76	0.77	3452
weighted avg	0.91	0.91	0.91	3452

This Repository works

-- Load saved models and run sample sentence --

- Testing the SystemA with an example sentence
([SystemADistilbert_Testing.ipynb](#)).

A

- Testing the SystemB with an example sentence
([SystemBDistilbert_Testing.ipynb](#)).

B

His	:	0
father	:	0
was	:	0
a	:	0
surveyor	:	0
and	:	0
tavern	:	0
owner	:	0
who	:	0
became	:	0
close	:	0
friends	:	0
with	:	0
William	:	B-PER
Henry	:	I-PER
Harrison	:	I-PER
while	:	0
the	:	0
two	:	0
served	:	0
together	:	0
in	:	0
the	:	0
War	:	B-EVE
of	:	I-EVE
1812	:	0
.	:	0

His	:	0
father	:	0
was	:	0
a	:	0
surveyor	:	0
and	:	0
tavern	:	0
owner	:	0
who	:	0
became	:	0
close	:	0
friends	:	0
with	:	0
William	:	B-PER
Henry	:	I-PER
Harrison	:	I-PER
while	:	0
the	:	0
two	:	0
served	:	0
together	:	0
in	:	0
the	:	0
War	:	0
of	:	0
1812	:	0
.	:	0

This Repository works

-- Findings -

- The testing set is missed a label ('I-BIO' - there is only 30 labels) compared to training set (31) and validation set (31).
 - Time for running:
 - SystemA: 5619.6813 s ~ 1.56 hour. (31 Labels)
 - SystemB: 5557,2894 ~ 1.5436915.
- The times for training of 2 systems are not far different. This is because there are changes in the last layer for the output layer with Transformers. Other layers are trained in similar ways. However, it can be different with larger datasets because I only test these systems on small random data from the original data because of time limitations and my machine.

This Repository works

-- Findings -

- Accuracy (Precision | recall | f1 score)
 - System A: 1.0 | 0.83 | 0.88
 - System B: 0.99 | 0.90 | 0.91

→ System B outperformed System A on average score. This is because of fewer labels compared to system A.
- Limitations:
 - Randomly small parts of the data can lead to incorrect evaluation of the models. Therefore, It needs to run with a powerful computer for all datasets.
 - Because of time limitations, I only tested on the default learning rate (example) from HuggingFace. It can change the learning rate with different optimizers like Adam. ([See more here](#))
 - Epoch = 10. It needs to be tested with more epochs. Note that I run 1 epoch with a full training set. It took more than 8 hours. According to the system, you can consider how many epochs you need for training. GPU is another solution to reduce the time for training.

Thanks

